

On the parallelization of a N-body simulation

ANDRÉ MENDES, Universidade do Minho, Portugal

BRUNO JARDIM, Universidade do Minho, Portugal

ROBERTO FIGUEIREDO, Universidade do Minho, Portugal

This report documents the parallelization techniques implemented on a N-body simulation. This was achieved through the OpenMP[1] API. After carefully analyzing and tested different optimization locations, the resulting program had a xxx% speedup compared to the original, sequential implementation. This result shows that physical, N-body simulations are an excellent target for parallel code optimizations. As even very simple optimization techniques are able to produce a huge speedup.

Additional Key Words and Phrases: Parallel, Parallelization, Code Optimization, Simulation, Physics Simulation, OpenMP

ACM Reference Format:

André Mendes, Bruno Jardim, and Roberto Figueiredo. 2026. On the parallelization of a N-body simulation. 1, 1 (January 2026), 8 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 Introduction

With the gradual decline of Moore's Law[2], the time of swapping an older chip for a newer one and experiencing exponential speedups is coming to an end. This, in a sense, limits the computing power a certain problem can have to still be tractable. That is, if we were dealing with a single computing core on a machine.

Fortunately, that is not the case. Nowadays, even the simplest computers have multiple computing cores. Unfortunately, this parallelism is often misused or even neglected, making the device, essentially, a single-core computer.

As such, this demands a narrative shift to parallelism as design principle. In fact, this is precisely the direction in which hardware is going, with more cores instead of faster ones, as we are reaching the physical limits of silicon. Thus, performance gains are not dependent on how fast a given CPU is, but how the parallelism of a given CPU can be exploited.

This issue becomes quite significant when it comes to N-body problems. These problems are notorious for their computational complexity, this is due to the pairwise calculations between every body in the system. As such, this requires a $O(N^2)$ time complexity, meaning that even small increases in particle count could dramatically increase the runtime of such an algorithm. In real-world scenarios, these applications require millions, if not billions, of particles in order to simulate any interesting phenomena. Thus making serial implementations of these applications impractical. This necessitates parallelization for these kinds of problems. However, parallelism isn't a solution without any drawbacks, there are a lot of considerations when parallelizing any algorithm. Such as the synchronization overheads, load imbalances, cache coherence effects, and even, memory bandwidth

Authors' Contact Information: André Mendes, andre@mail.com, Universidade do Minho, Braga, Portugal; Bruno Jardim, bruno.f.jardim@inesctec.pt, Universidade do Minho, Braga, Portugal; Roberto Figueiredo, roberto@mail.com, Universidade do Minho, Braga, Portugal.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2026 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM XXXX-XXXX/2026/1-ART

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

limitations. This article goes over the parallelization of a N-body simulation on multicore CPUs, focusing on memory behavior, scalability and algorithmic design choices.

2 Background

2.1 Overview of the N-body Problem

The N-body problem aims to simulate the time evolution of a given system with N particles, where all particles affect one another. Each time step necessitates the computation of the net forces that act on every single particle based on the states of all others and then updating their position and velocity. This results in an $O(N^2)$ complexity on every time step. Making the problem increasingly more expensive as N grows.

Due to this design the N-body problem is a well-suited candidate for parallelization. This is further reinforced by the fact that the force computations on individual particles are independent within that time step. However, most implementations on multicore CPUs are constrained by memory bottlenecks, cache coherence effects and even synchronization overheads. Thus, achieving linear speedups is not as trivial as it might seem.

Performance is heavily influenced by memory access patterns along with the in-memory representation of the given datastructures. As they can significantly improve cache utilization and throughput. All of these characteristics make the N-body problem a very useful benchmark for parallel performance on modern CPU architectures.

2.2 Parallel Computing on Multicore CPUs with OpenMP

OpenMP[1] is a commonly used programming model for shared-memory parallelism on multicore CPUs, providing a directive-based approach for defining parallel regions, work sharing, and synchronization. It allows existing sequential code to be parallelized incrementally, while maintaining a single shared address space. However, this abstraction does not remove the underlying costs associated with thread management, synchronization, and memory access, and as such performance still depends heavily on implementation details.

Parallel execution in OpenMP is most often achieved by distributing loop iterations or tasks across a team of threads. This maps well to data-parallel workloads such as N-body force calculations. However, scalability is influenced by scheduling choices such as the granularity of parallel work, and the amount of implicit synchronization introduced by constructs like barriers and reductions. Thus, suboptimal scheduling or excessive synchronization can limit performance improvements.

Memory behavior remains a central concern in OpenMP programs. Threads operate on cache-coherent shared memory, and frequent access to shared data can increase cache coherence traffic. This can lead to issues such as false sharing and reduced cache efficiency, which thus negatively affect scalability. As such, achieving good performance with OpenMP requires careful consideration of data layout, access patterns, and synchronization placement, in addition to exposing sufficient parallelism.

3 Sequential Implementation of the N-Body Simulation

The sequential implementation of the N-body simulation follows a straightforward time-stepping approach based on direct pairwise force computation. At each simulation step, the physics engine computes gravitational accelerations for all bodies and then updates their velocities and positions using explicit numerical integration. This implementation serves as the performance and correctness baseline against which parallel versions are evaluated.

The core computation is performed in the `integrateStep` function, which operates on a vector of bodies. For a system of N bodies, the function first allocates an auxiliary array to store the acceleration vectors for each body. The accelerations are initialized to zero at the beginning of each step, ensuring that force accumulation from the previous iteration does not persist. As such, force computation and state updates are cleanly separated.

Gravitational forces are computed using a double-nested loop that iterates over all pairs of bodies. For each body i , the contribution from every other body j is accumulated, excluding self-interactions. The force calculation is based on the classical inverse-square law, with a small softening term added to the distance computation to avoid numerical instability when bodies are very close. This approach results in a computational complexity of $O(N^2)$ per simulation step, which quickly becomes the dominant cost as the number of bodies increases.

Once all accelerations have been computed, a second loop updates the velocities and positions of each body using a simple explicit integration scheme. Velocities are updated based on the computed accelerations and the time step size, and positions are then advanced using the updated velocities. This integration method is computationally inexpensive but conditionally stable, and thus appropriate for short time steps and as a baseline reference.

Overall, the sequential implementation emphasizes clarity and correctness rather than performance optimization. It makes no attempt to reduce asymptotic complexity or exploit hardware parallelism. However, its simple structure and deterministic execution order make it well suited as a reference point for evaluating the effects of parallelization, memory behavior, and scalability in subsequent implementations.

4 ILP Optimizations

The first step when optimizing the sequential implementation is exploring the inherent capabilities of modern CPUs to execute multiple instructions concurrently. Even without introducing thread-level parallelism, the baseline N-body implementation exposes opportunities for improved instruction-level parallelism through careful ordering of arithmetic operations and memory accesses. As such, changes such as reducing control dependencies, increasing the number of independent operations, and improving data locality can significantly increase single-core performance. These optimizations preserve the sequential execution model while providing a more efficient foundation for subsequent shared-memory parallelization.

```
void PhysicsEngine :: integrateStep (SoABodies& bodies , double dt) {
    const double G = 6.67430e-11;
    size_t n = bodies.size();

    // Reset accelerations
    std::fill(bodies.acc_x.begin() , bodies.acc_x.end() , 0.0);
    std::fill(bodies.acc_y.begin() , bodies.acc_y.end() , 0.0);
    std::fill(bodies.acc_z.begin() , bodies.acc_z.end() , 0.0);

    for (size_t i = 0; i < n; ++i) {
        double ax = 0.0 , ay = 0.0 , az = 0.0;

        double px = bodies.pos_x[i];
        double py = bodies.pos_y[i];
        double pz = bodies.pos_z[i];

        for (size_t j = 0; j < n; ++j) {
            // if (i == j) continue; removed to optimize performance

            double rx = bodies.pos_x[j] - px;
```

```

double ry = bodies.pos_y[j] - py;
double rz = bodies.pos_z[j] - pz;

double distSq = rx * rx + ry * ry + rz * rz + 1e-9;
double inv_dist = 1.0 / std::sqrt(distSq);
double inv_dist3 = inv_dist * inv_dist * inv_dist;
double factor = G * bodies.mass[j] * inv_dist3;
ax += factor * rx;
ay += factor * ry;
az += factor * rz;

}

bodies.acc_x[i] += ax;
bodies.acc_y[i] += ay;
bodies.acc_z[i] += az;
}

for (size_t i = 0; i < n; ++i) {
bodies.vel_x[i] += bodies.acc_x[i] * dt;
bodies.vel_y[i] += bodies.acc_y[i] * dt;
bodies.vel_z[i] += bodies.acc_z[i] * dt;

bodies.pos_x[i] += bodies.vel_x[i] * dt;
bodies.pos_y[i] += bodies.vel_y[i] * dt;
bodies.pos_z[i] += bodies.vel_z[i] * dt;
}
}

```

The implementation exposes a significant amount of instruction-level parallelism through careful ordering of arithmetic operations and data access patterns. In the inner force computation loop, independent floating-point operations are structured in a way that allows the CPU to overlap execution and hide instruction latencies. In particular, the accumulation of acceleration components along each axis is performed using separate scalar variables, enabling the processor to execute these operations concurrently where possible.

Several code-level decisions further improve ILP. Particle position values for the outer-loop body are loaded once into local variables before entering the inner loop, reducing redundant memory accesses and allowing the compiler to better schedule dependent instructions. Additionally, the use of a structure-of-arrays data layout results in contiguous memory accesses, which improves cache behavior and allows the processor to issue multiple independent load instructions in parallel.

The removal of the conditional branch that excludes self-interactions eliminates a potential control dependency inside the inner loop. This reduces branch misprediction penalties and allows the instruction pipeline to remain fully utilized. As such, the inner loop consists primarily of straight-line code with predictable control flow, which is well suited for both static compiler scheduling and dynamic out-of-order execution. While these changes do not alter the algorithmic complexity, they improve single-core efficiency by increasing effective instruction throughput.

4.1 Single Instruction, Multiple Data (SIMD)

In addition to instruction-level parallelism, the implementation exposes opportunities for data-level parallelism through SIMD execution. The inner force computation loop operates on arrays of particle data and applies the same sequence of arithmetic operations to each element. This regular structure makes the loop a suitable candidate for vectorization, allowing multiple interactions to be processed simultaneously using wide vector registers.

The use of a structure-of-arrays data layout is particularly beneficial for SIMD execution. Particle positions, velocities, and masses are stored in separate contiguous arrays, enabling the compiler to generate vector load and store instructions with predictable memory access patterns. As such, multiple values of position or mass can be loaded into vector registers in a single instruction, improving throughput and reducing loop overhead.

SIMD parallelism is explicitly encouraged through the use of OpenMP `simd` directives on both the force accumulation loop and the state update loop. In the force computation, a reduction clause is used to safely accumulate partial results into scalar acceleration variables. This allows the compiler to vectorize the loop while preserving correctness, even though the accumulation introduces loop-carried dependencies at the scalar level. The removal of conditional branches within the loop further simplifies vectorization by maintaining a uniform control flow across iterations.

Similarly, the velocity and position update loop applies identical operations to independent elements of the state arrays and can be vectorized without requiring reductions. While SIMD execution does not change the algorithmic complexity of the simulation, it increases effective arithmetic throughput and improves single-core performance. As such, SIMD optimizations complement instruction-level parallelism and provide a more efficient foundation for thread-level parallelization in the shared-memory implementation.

```
void PhysicsEngine :: integrateStep (SoABodies& bodies , double dt) {
    const double G = 6.67430e-11;
    size_t n = bodies.size();

    // Reset accelerations
    std::fill(bodies.acc_x.begin() , bodies.acc_x.end() , 0.0);
    std::fill(bodies.acc_y.begin() , bodies.acc_y.end() , 0.0);
    std::fill(bodies.acc_z.begin() , bodies.acc_z.end() , 0.0);

    for (size_t i = 0; i < n; ++i) {
        double ax = 0.0 , ay = 0.0 , az = 0.0;

        double px = bodies.pos_x[i];
        double py = bodies.pos_y[i];
        double pz = bodies.pos_z[i];

        #pragma omp simd reduction(+:ax, ay, az)
        for (size_t j = 0; j < n; ++j) {
            // if (i == j) continue; removed to optimize performance

            double rx = bodies.pos_x[j] - px;
            double ry = bodies.pos_y[j] - py;
            double rz = bodies.pos_z[j] - pz;
```

```

double distSq = rx * rx + ry * ry + rz * rz + 1e-9;
double inv_dist = 1.0 / std::sqrt(distSq);
double inv_dist3 = inv_dist * inv_dist * inv_dist;
double factor = G * bodies.mass[j] * inv_dist3;
ax += factor * rx;
ay += factor * ry;
az += factor * rz;

}

bodies.acc_x[i] += ax;
bodies.acc_y[i] += ay;
bodies.acc_z[i] += az;
}

#pragma omp simd
for (size_t i = 0; i < n; ++i) {
    bodies.vel_x[i] += bodies.acc_x[i] * dt;
    bodies.vel_y[i] += bodies.acc_y[i] * dt;
    bodies.vel_z[i] += bodies.acc_z[i] * dt;

    bodies.pos_x[i] += bodies.vel_x[i] * dt;
    bodies.pos_y[i] += bodies.vel_y[i] * dt;
    bodies.pos_z[i] += bodies.vel_z[i] * dt;
}
}

```

While SIMD vectorization improves single-core performance by exploiting data-level parallelism within individual loops, it does not address parallelism across multiple cores. Thread-level parallelization and SIMD execution therefore target different levels of the hardware execution model. As such, SIMD optimizations can be applied independently of the chosen threading strategy and serve to increase the amount of useful work performed per core. This combination allows both forms of parallelism to complement one another, improving overall scalability when the simulation is executed on multicore CPUs.

5 Parallel Implementation

The parallel implementation builds on the sequential baseline by introducing shared-memory parallelism using OpenMP. Rather than restructuring the algorithm, parallelism is applied to the main phases of the time-stepping procedure: acceleration reset, force computation, and state update. These phases are executed within a single parallel region, allowing threads to be reused across multiple work-sharing constructs and reducing parallel region overhead. As such, the implementation combines thread-level parallelism across bodies with data-level parallelism inside the force computation.

```

void PhysicsEngine :: integrateStep (SoABodies& bodies, double dt) {
    const double G = 6.67430e-11;

```

```

size_t n = bodies.size();

#pragma omp parallel
{
    // Reset accelerations
#pragma omp for schedule(static)
    for (size_t i = 0; i < n; ++i) {
        bodies.acc_x[i] = 0.0;
        bodies.acc_y[i] = 0.0;
        bodies.acc_z[i] = 0.0;
    }

#pragma omp for schedule(static)
    for (size_t i = 0; i < n; ++i) {
        double ax = 0.0, ay = 0.0, az = 0.0;

        double px = bodies.pos_x[i];
        double py = bodies.pos_y[i];
        double pz = bodies.pos_z[i];

#pragma omp simd reduction(+:ax, ay, az)
        for (size_t j = 0; j < n; ++j) {
            // if (i == j) continue; removed to optimize performance

            double rx = bodies.pos_x[j] - px;
            double ry = bodies.pos_y[j] - py;
            double rz = bodies.pos_z[j] - pz;

            double distSq = rx * rx + ry * ry + rz * rz + 1e-9;
            double inv_dist = 1.0 / std::sqrt(distSq);
            double inv_dist3 = inv_dist * inv_dist * inv_dist;
            double factor = G * bodies.mass[j] * inv_dist3;
            ax += factor * rx;
            ay += factor * ry;
            az += factor * rz;

        }
        bodies.acc_x[i] += ax;
        bodies.acc_y[i] += ay;
        bodies.acc_z[i] += az;
    }
}

```

```

#pragma omp for schedule(static)
for (size_t i = 0; i < n; ++i) {
    bodies.vel_x[i] += bodies.acc_x[i] * dt;
    bodies.vel_y[i] += bodies.acc_y[i] * dt;
    bodies.vel_z[i] += bodies.acc_z[i] * dt;

    bodies.pos_x[i] += bodies.vel_x[i] * dt;
    bodies.pos_y[i] += bodies.vel_y[i] * dt;
    bodies.pos_z[i] += bodies.vel_z[i] * dt;
}
}
}

```

5.1 Parallelization Strategy

The primary source of parallelism is the outer loop of the force computation, where each iteration computes the total acceleration acting on a single body. This loop is parallelized using an OpenMP work-sharing construct, assigning distinct bodies to different threads. Since each iteration accumulates forces independently into local variables, no synchronization is required within the inner loop.

A single parallel region is used to cover all major computation phases within a time step. Inside this region, multiple `for` constructs distribute work for acceleration reset, force computation, and state updates. A static scheduling policy is employed, as the workload per iteration is uniform and predictable. This choice reduces scheduling overhead and provides good load balance without requiring dynamic scheduling or task-based parallelism.

5.2 Shared Memory Considerations

The implementation follows a shared-memory model in which all threads access the same particle data arrays. During the force computation, threads read shared position and mass data while writing only to private acceleration entries, which avoids data races and the need for synchronization. This access pattern is well suited to OpenMP-style parallelism, although scalability can still be influenced by memory bandwidth and cache coherence effects as thread count increases.

5.3 Memory and Cache Behavior

Performance is strongly affected by how particle data is accessed in memory. The structure-of-arrays layout improves spatial locality and enables more efficient cache usage during the force computation. Writing acceleration data to distinct array elements reduces false sharing, but cache line alignment and memory traffic remain relevant factors, particularly on systems with many cores.

References

- [1] Openmp application programming interface. <https://www.openmp.org>, 2026. Accessed 2026.
- [2] G.E. Moore. Cramming more components onto integrated circuits. *Proceedings of the IEEE*, 86(1):82–85, 1998.